

Delegating Work to AI

A practical test for deciding what to hand off, what to review,
and what humans still own

Philip Budden

Specify

What are we asking for?

Verify

Can we check it?

Judge

What stays human?

Own

Who stands behind it?

About Me

Phil Budden

Senior Data Engineer & Data Platform Lead, Codi Group

- Working across **data, AI strategy and responsible AI adoption** in social housing
- Technical practitioner increasingly using AI to **delegate substantial knowledge and engineering work**
- Interested in the gap between **what AI can produce** and **what we can dependably hand over to it**

The Practical Delegation Test grew out of a practical question: how much of this work should I actually hand over?

The Productivity Story Is Too Narrow

AI can generate plausible work quickly, but generation is only one stage in the workflow.

Define

intent, audience,
constraints

Generate

now much cheaper

Verify

still needs
attention

Integrate

fit with the real
workflow

Own

accountability
remains

The Better Question

TOO NARROW

Can AI produce this?

MORE USEFUL

Can this human-and-AI system deliver this dependably?

Count specification, verification, judgment and accountability.

Review Debt

AI can create candidate work faster than qualified people can check it.

Generate

candidate work arrives quickly

Queue

unchecked work accumulates


Review

expert attention remains finite

Where The Test Sits

Where the Practical Delegation Test Sits

Task-level AI decisions should operate inside wider obligations and controls, not outside them.



The Practical Delegation Test

01

Can I describe the outcome clearly?

success, failure, audience, constraints

02

Can I verify the result cheaply enough?

time, expertise, detectability, reversibility

03

What judgment must remain human?

belief, risk, commitments, trade-offs

04

Where does the uncertainty go?

absorbed, named, escalated, or dumped

Question 1: Outcome Clarity

Audience

Who is the output for?

Acceptable

What would make this good enough?

Unusable

What would make it unsafe or wrong?

Constraints

What must it include, exclude, or obey?

If the brief is vague, use AI for exploration rather than execution.

Outcome Clarity Failure Pattern

WEAK TASK

Write something good about this.

->

STRONGER TASK

Draft a two-page decision brief for this audience, covering these points, excluding these claims, and naming remaining uncertainty.

Weak clarity moves the task toward exploration.

Question 2: Verification Cost

Time

How long does meaningful review take?

Expertise

Who can detect important failures?

Detectability

Are errors visible or hidden?

Reversibility

Can missed errors be corrected?

Check structure

Is there a repeatable test?

Verification Is Not One Thing

Cheap to check

- Narrow tests
- Reconciliations
- Known examples
- Objective criteria

Expensive to check

- Subtle omissions
- Tacit context
- Live consequences
- Persuasive but untested claims

Question 3: Retained Human Judgment

Draft

AI can produce candidate
wording

Compare

AI can surface options

Commit


humans own promises

Accept risk


humans own trade-offs

`Human in the loop` is too vague.

Judgment Is Not Always One Person



Question 4: Where Does The Uncertainty Go?



Good handoffs absorb, name, or escalate uncertainty instead of passing it on as polish.

Legitimate Review Vs Hidden Transfer

Legitimate review

- Expected, scoped and resourced
- Residual uncertainty named
- Reviewer can challenge

Hidden transfer

- Implicit, improvised or dumped
- Uncertainty obscured by polish
- Reviewer pressured to accept

From Questions To Delegation Levels

Direct delegation

clear, testable, low
residual judgment

Checkpointed delegation

useful execution,
important review points

Human-led with AI support

ambiguity or judgment
remains central


Do not delegate

too unclear, sensitive or
hard to verify

Three Modes Of AI Use

Three Modes of AI Use

Modes describe the kind of use; delegation levels describe how much authority and review burden is being delegated.



**Do not mistake assistive use
for safe executional delegation**

Worked Examples

Meeting summary

clear enough, but
decisions, actions and
tone need review

Checkpointed

Unit tests

narrow task, objective
execution path, quick
failure signals

Direct + review

Complaint response

format can be supported,
but remedy and
commitments stay
human


Human-led

Redundancy ranking

severe consequences,
contestable assumptions,
legitimacy dominates

Do not delegate

Agentic Workflow Example



Apply the test at workflow or sub-workflow boundaries.

Final review may hide routing, retrieval, classification and escalation failures.

Accountability Controls

Approval

explicit release
points

Ownership

named
accountable
person

Evidence

what was checked

Challenge

reviewers can
reject

Traceability

proportionate
record

Review Power Matters

Time

enough space to check

Context

enough information to
judge

Authority

permission to reject

Safety


cover to challenge speed
pressure

A review step without power is often just a rubber stamp with nicer paperwork.

Maturity Over Time

Delegation Can Mature Over Time

A task may become safer to delegate as specification, verification, and evidence improve.



A Lightweight Team Worksheet

Task

name the work and
audience

Checks

verification path and
reviewer

Judgment

what stays human

Decision

level, uncertainty,
controls

The worksheet should fit in a meeting note. If it becomes ceremony, use a team default or split the task.

What Good Use Looks Like

Cleaner handoffs

less work arriving as polished uncertainty

Faster safe delegation

stable low-risk tasks move more easily

Better disagreement

borderline tasks expose the real issue

Less cleanup

review debt is noticed before it compounds

The Leadership Ask

Do not govern AI only through tool approval.

Workflow

where the work happens

Review debt

where checking
accumulates

Uncertainty

where doubt is carried

Accountability

who can answer for it

Closing

NOT THE GOAL

Maximum delegation

THE GOAL

**Dependable work with
clear ownership of
remaining judgment**

Use the four questions before cheap generation creates expensive ambiguity.

Q&A

What kind of task are you currently unsure whether to delegate?

01

Can I describe the outcome clearly?

02

Can I verify the result cheaply enough?

03

What judgment must remain human?

04

Where does the uncertainty go?

Find out more

philipbudden.co.uk/practical-delegation-test/

- **Canonical whitepaper** for senior leaders, sponsors and accountable owners
- **Practitioner guide** for applying the test to your own work
- **Feedback form**, with an option to respond anonymously

Please share successes, difficulties, disagreements with the framework or its questions, and work the four questions did not fit.